

Project 4: Financial Engineering from a Statistical Physics Approach

Michael Saybolt

5/7/2017

Contents

1	Introduction	2
2	Technical Details	3
2.1	General Considerations	3
2.2	Parameter Choice	3
2.3	Termination Conditions	4
3	Big Data Challenges	4
3.1	Dealing with Large Sets of Data	4
3.2	Computation of Large Sets of Data	6
3.2.1	Parallelization Attempt	6
3.2.2	Crude but Effective Method	6
4	Results	6
4.1	Default	6
4.2	Savings	8
4.3	Similar Wealth	9
4.4	Adding Previous Interactions	9
5	Conclusions and Future Work	13

Abstract

This project explores the generation and simulation of an economic system. Financial transactions among agents are simulated using Monte Carlo methods. The resulting distribution of wealth is then studied. Various parameters can be tuned to create a system that accurately models real wealth distributions from various countries.

1 Introduction

The end goal is to analyze the distribution of wealth as a function of agents' money/income, m . The result should follow a Pareto distribution [1]. In addition, the model was made more realistic by following the work of [Patriarca and collaborators](#), adding a quantity λ responsible for making the agents save a certain fraction of their money during a transaction [2]. Following the work of [Goswami and Sen](#), the agents can be further enhanced to model

tendencies to interact with other agents of similar income, as well as previous interactions[3].

2 Technical Details

2.1 General Considerations

In all levels of modeling, there are some basic considerations that may differ slightly from the papers. In general, the distribution of money follows

$$w_m \propto m^{-1-\alpha},$$

with $\alpha \in [1, 2]$ [1].

A certain number of agents N are created for each Monte Carlo iteration, and they are randomly selected to perform transactions. Unlike the paper, they are each assigned some equal start money, m_0 , instead of a random value from a uniform distribution. The choosing of the agents, however, is still done using a uniform distribution. For some number of Monte Carlo cycles, for some number of transactions, a pair of agents (i, j) are chosen by the uniform distribution, and exchange money. Money is conserved, thus

$$m_i + m_j = m'_i + m'_j. \quad (1)$$

The random number ϵ , generated from a uniform distribution, is used to determine the fraction of money that is changed. The money of agents' i and j become

$$m'_i = \epsilon(m_i + m_j),$$

leading to

$$m'_j = (1 - \epsilon)(m_i + m_j).$$

No agents can go into debt (only finitely smaller transactions), and due to the conservation law, the system relaxes towards a Gibbs distribution.

2.2 Parameter Choice

Parameters are chosen initially such that when the system reaches its termination condition, there are not 'holes' in the data set. Without sufficient interactions, and more importantly enough Monte Carlo cycles to even out

the dataset, there will be parts of the distribution of wealth with no agents in that section. Taking the log of the distribution of wealth will look uneven and be inaccurate as a result, when it should look relatively linear for the most basic economic model (no money saved or other caveats to interaction). Additionally, sufficient start money seemed to play a role (even though it seems it shouldn't, as all floats tested can be divided or multiplied many times themselves without significant accuracy loss), however there was not enough CPU time to verify this trend.

2.3 Termination Conditions


The termination condition is met when the system has reached some sort of equilibrium. This can be guessed by looking at the data and seeing when it quits changing much with additional transactions, or for a more quantitative approach, the variance amongst the agents can be observed as the transactions go on, which will stabilize when the system reaches equilibrium. Figures 1a, 1b, and 1c show the variance against number of transactions and 10 Monte Carlo cycles for clarity. Choosing `num_transactions` to ensure the system is at equilibrium is no longer a guessing game. This tends to be very predictable, and unless the number of agents `N` is changed, the system's convergence to equilibrium is easily predicted within an order of magnitude.

3 Big Data Challenges

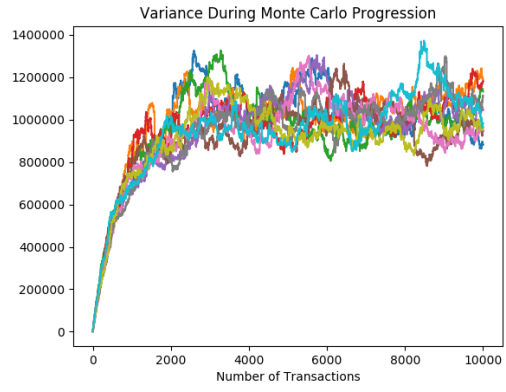
3.1 Dealing with Large Sets of Data

Unlike previous projects, Monte Carlo simulations involve a lot of CPU time to generate a sufficient pool of pseudo randomly generated data. While some general trends can be observed with less iterations for debugging, to truly check the details, sufficient numbers of experiments must be done.

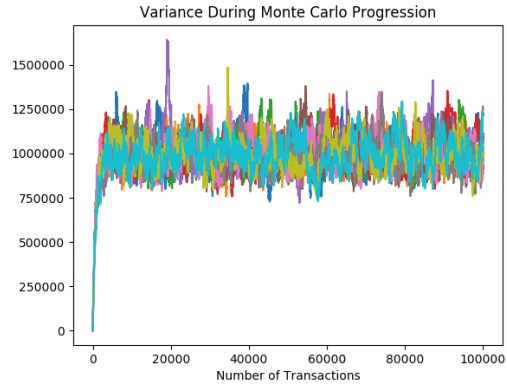
This prompted the need for modularizing the code, building a library for commonly used functions, adding load and store functions, as well as a separate program `quikplot` to just plot data. Running through all that computation just to plot something and then discard it after making changes to the code is incredibly wasteful and inefficient, so after the cycles complete, it optionally writes to a file and/or plots, using functions contained in a shared financial library. This way, a separate plotting function can load the



(a) $\text{num_transactions}=10^3$



(b) $\text{num_transactions}=10^4$



(c) $\text{num_transactions}=10^5$

Figure 1: Variance of 10 MC cycles with $N = 500$

files into variables and call the exact same plotting function to recreate and display that environment.

3.2 Computation of Large Sets of Data

3.2.1 Parallelization Attempt

When dealing with such a CPU intensive task, parallelization immediately comes to mind. Unfortunately parallelizing the Monte Carlo cycles did not seem trivial, at least in the stage of the code that was first attempted to be parallelized. This seemed strange, as individual cycles do not depend on each other unless the termination condition relies on comparing other cycles and must index them numerically during computation.

3.2.2 Crude but Effective Method

Since CPU time is of essence near the end of the semester, utilizing as much of it as possible is still a top priority. While my 4 GHz antenna simulation targeted desktop can kick out more flops than the engineering building's compute servers core-per-core, the compute servers have many, many more cores. Just dying for a parallel implementation of this code. Since parameters needed to be varied for the extra details and caveats of the transactions, multiple python instances can be started separately, and then data combined at the end, programatically, but not truly 'parallel' code.

Of course care was taken to use the least loaded one, but running 30+ instances of Python on cores half as fast was still a much faster way to acquire large sets of reasonably accurate data!

4 Results

A formal discussion of the results of adjusting the model to more accurately represent real economic principals is presented.

4.1 Default

The 'default' model simply follows exponential decay as mentioned in the introduction [1]. Agents do not save any money, nor do they discriminate




Figure 2: Taking over the egr compute servers




Figure 3: The 'default' Gibbs distribution, $N=500$, $\text{num_transactions}=10^7$, $\text{num_experiments}=10^2$

based on wealth or previous transaction. There are many agents with little money, and very few with a lot of money.

Taking the log of this function results in a relatively straight line, so long as enough interactions take place. Despite the high number of transactions, the lower number of experiments results in poor convergence despite taking just as many flops as if $\text{num_transactions}=10^6$ and $\text{num_experiments}=10^3$.

4.2 Savings

To add realism or study a hypothetical economic systme, agents can save a fraction of their money, λ , in a transaction. While equation (1) for conservation of money still holds, the new money of the agents is modeled as follows

$$m'_i = \lambda m_i + \epsilon(1 - \lambda)(m_i + m_j),$$

and

$$m'_j = \lambda m_j + (1 - \epsilon)(1 - \lambda)(m_i + m_j),$$

When agents save money, a strong 'middle class' is born. The more they save, the stronger this gets, and less agents have little to no money to deal with. The distribution goes from resembling a Gibbs distribution, to almost

a perfect Gaussian. Results for saving a quarter, half and 90% of their money is shown by figures

wtf 4a, 4b, and 4c respectively.

Note the importance of having enough MC cycles to get an accurate representation of the correct distribution. Figure 4d shows the last MC cycle in this data set alone without 'averaging' them all together; the data is much more sporadic.

4.3 Similar Wealth

Agents can also be tailored to be more likely to interact with other agents of similar wealth, or geographic proximity, regardless of how much they save. This seems to help the lower and mid-income agents, and if the log is taken, somewhat resembles the plots in Goswami and Sen[3]. It seems there may be a bug since increasing α does not have as much effect on the curvature as expected.


Taking it to the max to really see the influence, there is a difference with $\alpha=10$, but it is not as much change as expected.

The figures are enlarged to show the x-axis scale, which does change considerably. The effect of larger α causing the agents to stick together results in a lower upper bound for maximum money.

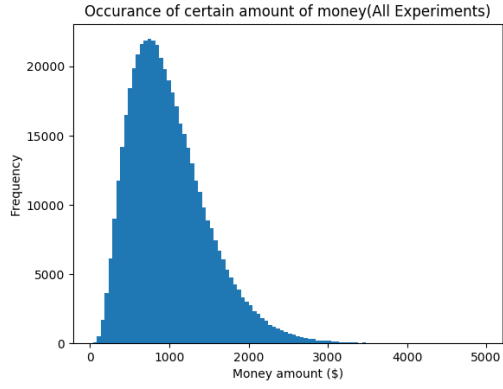
4.4 Adding Previous Interactions

People may have a tendency to want to continue to do business with previous clients, assuming it was a good transaction. Multiplying the probability of an interaction based on similar wealth, with an additional component containing a comparison to the maximum transactions between any agents, a confidence factor can be added in, increasing the likelihood of a transaction occurring the more previous transactions some agents have with each other.


Unfortunately, the result looks like it has been skewed by a bug where the probabilities are not scaled to each other properly. You see one effect clearly overtake the other. The small shape of the curvature change from $\alpha=1$ to $\alpha=2$ is present, delaying the effect of the previous transactions until higher γ . Eventually, the distribution is taken over by clumping as previous transaction preference becomes the dominant factor. For some reason, a certain percentage of agents appear to be holding out and not interacting at all, even with 10^6 transactoins per experiment.




(a) $\lambda=0.25$, $N=500$



(b) $\lambda=0.5$, $N=500$



(c) $\lambda=0.9$, $N=500$



(d) $\lambda=0.9$, $N=500$, last MC cycle only

Figure 4: Distribution of wealth for `num_transactions`= 10^5 , `num_experiments`= 10^3




Figure 5: Distribution of wealth for $\alpha=0.5$




Figure 6: Distribution of wealth for $\alpha=1.0$




Figure 7: Distribution of wealth for $\alpha=1.5$




Figure 8: Distribution of wealth for $\alpha=2.0$




Figure 9: Distribution of wealth for $\alpha=10$

The Pareto distribution from Goswami and Sen [3] can be reproduced by plotting the top edges of these trials on a log/log scale as shown in figure 12

5 Conclusions and Future Work

Monte Carlo methods were used to model an economic system in similar ways as the works of Patriarca[2] and Goswami and Sen[3]. While the results turned out similarly to works cited in most cases, there were still potentially some strange bugs that were not completely eradicated.

Additionally, there were convergence issues in regards to the system reaching equilibrium. Perhaps more MC cycles would have fixed this, but to go up another order of magnitude would have at least doubled CPU time that just wasn't present. While crudely parallelizing the code to generate more results in a shorter time period certainly helped (else the log results would barely look linear), there were still some discrepancies.


Lastly, this model could include more details such as adding a system for taxation, or subdividing it up into smaller economies and examining effects of

exchange rates around the world, as well as removing the debt restriction so that agents can obtain debt, or sub-economic systems can become unstable and collapse, which unfortunately frequently happens with real people.


All in all, this was an interesting project and provided additional variety of challenges to solve. Implementing better methods for folder naming conventions and really more robust loading and saving of data would have been useful to set up in the very beginning. I just wish I had more time to let this run and really examine things. Modularized code has become increasingly important with the progression of these projects, and almost becomes its own optimization problem on determining what parts are worth the effort to make many iterations of to test and ensure the modular version works just as well and hopefully better. Many lessons were learned, and I leave with an enhanced programming environment and mental set of tools to assist with other projects.

References


- [1] V. Pareto, Cours d'economie politique, Lausanne, 1897.
- [2] M. Patriarca, A. Chakraborti, K. Kaski, Physica A **340**, 334 (2004).
- [3] S. Goswami and P. Sen, Physica A **415**, 514 (2014).




(a) $\alpha=1, \gamma=0$




(b) $\alpha=1, \gamma=1$




(c) $\alpha=1, \gamma=2$



(d) $\alpha=1, \gamma=3$




(e) $\alpha=1, \gamma=4$




(f) $\alpha=1, \gamma=5$


Figure 10: Distribution of wealth for $\alpha=1$, $N=500$, $\text{num_transactions}=10^6$, $\text{num_experiments}=10^3$




(a) $\alpha=2, \gamma=0$




(b) $\alpha=2, \gamma=1$




(c) $\alpha=2, \gamma=2$



(d) $\alpha=2, \gamma=3$



(e) $\alpha=2, \gamma=4$



(f) $\alpha=2, \gamma=5$

Figure 11: Distribution of wealth for $\alpha=2$, $N=500$, $\text{num_transactions}=10^6$, $\text{num_experiments}=10^3$




Figure 12: Pareto distribution from Goswami and Sen recreation